

Design Justification Report: A Tiled 128-MAC INT8 Convolution-Reduction Chiplet

ECE 410/510 HW4AI – Milestone 4

Kay Hynes

June 7, 2026

1. Problem and motivation

This project accelerates the multiply-accumulate reduction at the center of a representative 3x3 INT8 convolution from YOLO-nano. The selected layer has a 52x52x64 input, 128 output channels, and a reduction length of $3 \times 3 \times 64 = 576$ per output pixel and channel: 199,360,512 MACs (398,721,024 FLOPs).

In M1 profiling, `conv3x3_int8_vectorized` consumed 2.275 of 2.379 seconds across 15 calls (95.6%). The NumPy `im2col-plus-GEMM` implementation ran the layer in a median 160.9 ms at 2.478 GFLOP/s. This is the M4 comparison point.

The chiplet accelerates the reduction kernel only. A host supplies activation elements, loads weights, accumulates returned tile partials, and handles sliding-window generation, padding, and output conversion. This partition is what makes the design synthesizable, interface-connected, and measurable against the M1 baseline.

2. Roofline analysis

The M1 **algorithmic arithmetic intensity** counts the input feature map, weights, and final output once:

$$\begin{aligned} &398,721,024 \text{ FLOPs} / (173,056 + 73,728 + 346,112 \text{ bytes}) \\ &= 672.497 \text{ FLOP/byte} \end{aligned}$$

At this AI the convolution is compute-bound on both the M1 Pro reference and the originally proposed accelerator. This motivated a parallel INT8 MAC array: at this intensity, multiplier count is the lever, not external bandwidth.

The implemented chiplet does not achieve ideal reuse. Weight banks hold 64 entries per lane, so a 576-element reduction needs nine tiles, and each tile returns 128 INT32 channel partials over a 64-bit serialized stream. Per layer the transaction schedule is 73,728 weight-load beats, 1,557,504 compute beats, and 3,115,008 output beats — 37,969,920 stream bytes at 8 bytes/beat — for an **implemented-interface arithmetic intensity of 10.501 FLOP/byte**.

The kernel is compute-bound under ideal on-chip reuse; the implemented chiplet is serialization-bound. At the projected 114.536 MHz one 64-bit stream direction is rated 0.916 GB/s and the

schedule needs 0.889 GB/s. Figure 1 shows both AI points and the projected production point.




Figure 1: Final roofline. The M1 algorithmic point assumes ideal reuse at 672.5 FLOP/byte. The synthesized tiled chiplet operates at 10.5 FLOP/byte because it transports tile commands and serialized partial sums; combining measured cycles with the setup-limited post-CTS projection gives 9.335 GFLOP/s.

3. Precision and data format

The datapath uses signed INT8 activations and weights with signed INT32 accumulation. Each INT8 multiplication produces a signed 16-bit product, which is sign-extended before accumulation. INT32 is sufficient for a full 576-element reduction: even the conservative magnitude bound $576 \times 16,384$ is below 10 million, well within signed 32-bit range.

INT8 was selected for area-per-multiplier on an edge target. The precision study ran 1,000 deterministic 576-tap reductions against an FP32 reference, matching the full reduction reconstructed from nine hardware tiles: mean absolute error 0.0360, RMS error 0.0448, max 0.1503, relative MAE 0.564%. Sufficient for the reduction kernel; task-level YOLO accuracy needs a trained quantized model and is out of scope.

Each lane represents its running partial sum in carry-save form as two 32-bit registers. The inner loop computes a three-input XOR and a majority function bitwise instead of propagating a carry through a 32-bit adder every cycle. A single carry-propagate add resolves the result on the tile's **last** element. Carry-save is bit-exact relative to ordinary INT32 addition for the tested reductions.

4. Dataflow and architecture

The compute dataflow is **output-stationary and weight-streaming within each tile**. One activation is broadcast to all 128 lanes each compute cycle. Lane i receives the weight for output channel i and keeps that channel’s partial sum resident for the 64-element tile. At the end of the tile, all 128 channel partials are captured for serialization. This dataflow fits convolution reduction because all output channels consume the same activation element while using different weights.

The `mac_array` compute engine contains 128 parallel lanes and a three-stage pipeline: Stage A captures activation and weights, Stage B performs the signed 8×8 multiplication, and Stage C updates or resolves the carry-save accumulator. Shared `valid`, `first`, and `last` tags travel with the data. During a tile, the array accepts one reduction element per cycle and performs 128 useful MACs per compute cycle.

The synthesized production boundary is `accel_top`, shown in Figure 2. It adds four necessary chiplet functions:

1. A 64-bit opcode-tagged AXI4-Stream input.
2. 128 private 64-entry-by-8-bit inferred weight banks.
3. A registered activation/control broadcast stage.
4. A result buffer and serializer that returns one channel partial per beat.

The final scope is a **64-element tiled reduction chiplet**. Nine invocations reconstruct the full 576-element reduction; the host adds each returned INT32 partial inline as its serialized beat arrives, which the testbench models (no extra cycles beyond the drain stream). Sliding-window generation, padding, line buffers, and output quantization remain with the host. This is narrower than the M1 target but is the design represented consistently across final simulation, synthesis, and benchmark.

Figure 2. Final production architecture: host-managed nine-tile reduction through synthesized `accel_top`




Figure 2: Production architecture. The host loads one 64-weight tile per lane, streams 64 activation elements per output pixel, receives 128 serialized partials, and accumulates nine tiles for the full reduction.

Figure 3. Tiled output-stationary dataflow: 64-element chiplet partials, nine-way host accumulation




Figure 3: Output-stationary tile dataflow. Partial sums remain in each lane for one 64-element tile; nine returned tile partials are accumulated by the host.

5. Hardware interface

The chiplet exposes 64-bit AXI4-Stream input and output (TVALID/TREADY/ TDATA). AXI4-Stream is a published AMBA streaming protocol that connects naturally to host DMA or control logic on the assumed FPGA-SoC, with no chiplet-side change.

Input opcode 0x01 loads one signed INT8 weight into a selected lane and address. Opcode 0x02 carries a broadcast activation plus **first/last** tags. The output returns a 7-bit channel index and one signed INT32 partial per beat. The testbench drives complete weight, compute, and response transactions exclusively through these ports.

The interface is functionally correct but is the dominant performance limit. The serializer needs 128 output beats after every 64 compute beats, and **accel_top** has one result buffer (input TREADY drops while draining), so each tile must drain before the next can issue. The result path therefore costs more cycles than the compute path.

M1's bandwidth conclusion assumed a full 576-entry on-chip weight store and ideal feature-map reuse; that design would not be interface-bound. The implemented 64-entry tiled design is, and the benchmark reports it.

6. Verification

The final testbench **tb/tb_top.sv** instantiates **accel_top** at the synthesized **L_MAX=64** and drives only the production AXI4-Stream ports. It builds deterministic signed INT8 activations and weights for eight pixels and computes two independent references: every 64-element tile partial and every full 576-element result. Pixel 0, channel 0 uses alternating 127/-128 operands to exercise full-scale INT8 and large-magnitude carry-save accumulation; the other vectors provide varied signed values.

For each of nine tiles the testbench loads 128 x 64 weights, streams each pixel's 64 activation

elements, receives 128 serialized channel partials, and adds them into a host-side accumulator. It checks 9,216 tile partials and 1,024 reconstructed full results. The committed log reports:

```
partial_errors=0 full_errors=0
total_cycles=87988
backpressure_cycles=3 backpressure_errors=0 unstalled_schedule_cycles=87985
useful_macs_per_total_cycle=6.703
PASS: final synthesized-config accel_top tiled 576-element reduction
```

The test deliberately deasserts output **TREADY** for three cycles and still passes every ordering and value check. The 87,985-cycle unstalled schedule, not the injected stall, is used for layer extrapolation. The 6.7 MAC/cycle figure in this short run amortizes 73,728 weight beats over 8 pixels; the full-layer benchmark amortizes them over 2,704. The test validates the four cycle categories the benchmark extrapolates from: weight loads, compute beats, serialized drains, and protocol gaps. Figure 4 shows a compute tile entering the array and its partial result entering the serializer.




Figure 4: Annotated production waveform showing a 64-tap tile entering `accel_top` and its 128-channel partial result beginning serialized output.

7. Synthesis results

Synthesis: OpenLane 2 v2.3.10, sky130A / sky130_fd_sc_hd. The final configuration `synth/config.json` is `accel_top`, 128 lanes, L_MAX=64, 6.0 ns target, 3.3 mm x 3.3 mm

die.

The wrapper completed synthesis, floorplanning, placement, CTS, post-CTS timing repair, and global routing. The deepest snapshot contains 558,792 standard cells and 5.015 mm² of cell area. Dominant contributors: inferred weight-register banks, 128 multipliers and carry-save accumulators, clock tree, timing-repair buffers. The 65,536 weight bits alone account for at least 77.17% of the design’s 84,927 sequential cells; the 70,646 timing-repair buffers are 12.64% of all post-CTS cells. Global routing reached zero overflow on every metal layer.

At the post-CTS typical corner, the 6.0 ns target has setup WNS -2.730896 ns — a setup-limited projected period of 8.730896 ns, or **114.536 MHz**. The same snapshot has -1.852060 ns hold WNS and 32,070 reported hold violations, but register-to-register hold worst slack is +0.197419 ns with zero violations. There are two max-slew and one max-capacitance violations. Detailed routing and final timing repair did not complete, so these are unresolved; 114.536 MHz is a setup-limited estimate, not sign-off, and no slow-corner wrapper frequency is claimed.

The worst setup path starts at a result-path flip-flop, runs through fanout buffers and serializer mux/output logic, and ends at `m_tdata[26]` (-2.730896 ns slack). The interface, not the per-lane multiplier, is the wrapper timing limiter. The run used OpenLane’s generic fallback SDC because no `PNR_SDC_FILE` or `SIGNOFF_SDC_FILE` was supplied. A follow-on run needs an explicit host-interface SDC, completed detailed routing, and rerun setup/hold repair before interface timing closure is claimed.

Post-CTS power estimate: **1.018 W** at the typical corner with default switching activity, dominated by sequential logic and clock distribution. Pre-detailed-route and not workload-annotated. Analyzed at the 6.0 ns target rather than the slower setup-limited frequency. Used as an approximate energy result only.

8. Benchmark results

The final benchmark derives full-layer cycles from the transaction schedule measured by `tb_top.sv`. One layer requires:

```
73,728 weight-load cycles
1,557,504 compute-input cycles
3,115,008 serialized-output cycles
146,017 pipeline and protocol-gap cycles
= 4,892,257 total cycles
```

40.750 useful MACs per total chiplet cycle. At the projected 114.536 MHz this is 42.714 ms and **9.335 GFLOP/s** — a **3.77x projected speedup** over the M1 median (160.9 ms, 2.478 GFLOP/s) for the same FLOP count. M1 includes im2col/padding/output and M4 excludes host orchestration; the wrapper does not close timing. So this is a chiplet-schedule projection, not a demonstrated end-to-end operating point.

Power estimate x projected runtime gives 43.483 mJ/layer and 9.17 GFLOP/s/W. Arithmetic estimate only. No software energy ratio is claimed because the M1 CPU baseline did not include a measured power value.

M1 targeted 64 GFLOP/s at 250 MHz. The gap is the lower wrapper frequency projection, nine-way tiling forced by 64-entry weight banks, and serialized tile-partial output. Figure 1 reflects

these.

9. What did not work

The first major failure was scope: a full 576-entry weight bank for every lane was too large when inferred as flip-flops in the available educational flow. The final synthesizable production configuration therefore uses 64 entries and host-managed tiling. A better implementation would integrate SRAM macros so the complete 576-element weights remain on chip.

The second failure was treating the compute-array rate as accelerator throughput. The wide unit-test wrapper and the inner compute phase can sustain nearly 128 MACs/cycle, but the production wrapper must serialize 128 channel partials and has only one result buffer. When the final testbench was changed to 64-tap tiles, issuing pixels without waiting for serialization overwrote the result buffer. The corrected host protocol waits for each result drain, and the benchmark includes that cost. A next version should use double buffering, a FIFO, or a wider output stream to overlap compute and drain.

The attempted eight-bank broadcast optimization also did not work as intended. `accel_top` creates eight registered activation/control copies, but `mac_array` exposes only one scalar shared activation/control input, so only bank 0 is connected and the unused copies are eligible for synthesis removal. The final design therefore has registered broadcast staging, not a true eight-bank fan-out tree. A real banking fix requires splitting `mac_array` into banked instances or changing its interface so each lane group consumes its own registered copy.

The third major failure was timing closure. In the separate routed-lane experiment, carry-save accumulation moved the critical path away from the 32-bit accumulation adder and exposed the 8x8 signed multiplier as the new lane limiter. The full wrapper was slower still because of its weight storage, shared broadcast network, clock distribution, and routing. It also retains unresolved non-register-to-register hold, slew, and capacitance violations at the final snapshot. Pipelining the multiplier would improve the lane, while completed timing repair and hierarchical physical design would better address the full array. I would also provide an explicit host-interface SDC instead of the generic fallback so input/output timing and the serializer critical path are constrained against the assumed FPGA-SoC.

The fourth failure was full detailed routing. The bare array first failed because exposing all wide weights and results created 5,134 top-level pins. `accel_top` reduced that to 137 pins and completed global routing with zero overflow, but TritonRoute failed during detailed-routing track assignment on the available 16 GB host. The committed `openlane_run.log` records this failure, and the report uses only the deepest available full-wrapper post-CTS results.

Finally, power estimation is not workload-annotated. The 1.018 W result uses default activity and lacks final routed parasitics. A stronger follow-on would complete detailed routing and replay the production testbench with VCD or SAIF activity.

The design is narrower and slower than M1's target but is coherent: the same 64-tap chiplet is verified, synthesized, analyzed, and benchmarked, and reconstructs the 576-element reduction. The projected kernel advantage is real; timing closure and an end-to-end host benchmark would be needed to call it demonstrated.